Exam completed on: September 9, 2021

Score: 89% · 32/36

Exam questions

Show all | Show wrong only

✗ Question 1 of 36

Select the regular expression that returns the full year 2018, and not the single numbers individually.

- `[0-9]`
- ✗ `[2-8]+`
- `[0-9]*`
- ✓ `[0-9]+`

✓ Question 2 of 36

Select the scenario that would not make a feature a prime candidate for transformation.

- **few outliers**
- **dramatic skew**

long tail

✓ short tail

✓ Question 3 of 36

You believe that spam text messages are longer than real ones. What is one way to determine if you are correct?

Create a feature that contains bins for the length of each text message and plot the data on a pie chart.

✓ Create a feature that contains bins for the length of each text message and plot the data on a histogram.

Create a feature that contains bins for the length of each text message and plot the data on a scatter plot with a trend line.

✓ Question 4 of 36

After reading the data using pd.read_csv() into x, how can you tell the number of rows?

✓ len(x)

x.rows()

x.count()

rows(x)

✓ Question 5 of 36

Which step is important when creating a new feature?

✓ analyzing the dataset and coming up with a hypothesis

determining which lambda function to use

making sure that you always start with the preprocessed dataset

✗ Question 6 of 36

Why would you need to apply a transformation to your data?

✓ to lessen the effect of outliers to make better correlations

✗ to allow feature engineering to be applied to the data

to allow more creativity when using feature engineering methods

✓ Question 7 of 36

Which statement about training a gradient boosting model is TRUE?

Changing the `learning rate` parameter will have a very small effect on the overall results.

✓ An `n_jobs` parameter is not needed.

All decision trees will have unlimited depth.

✓ Question 8 of 36

How does gradient boosting work?

It forms a strong learner by discarding mistakes from prior test iterations.

It generates a range of input conditions for a machine learning model to test with.

✓ It combines weak learners together to form a strong learner by improving on mistakes from prior test iterations.

It generates boundary conditions for a machine learning model to test with.

✓ Question 9 of 36

Which two concepts are combined to create a more powerful tool to tune and evaluate machine learning models?

cross validation and transformation

cross validation and vectorizing

✓ grid search and cross validation

grid search and transformation

✓ Question 10 of 36

How would you define a stop word?

a frequently used word that appears before a period

a frequently used word that appears at the end of a sentence

✓ a frequently used word that doesn't contribute to the meaning of the sentence

a frequently used word that contributes to the meaning of the sentence

✗ Question 11 of 36

What is the purpose of a grid search?

to determine which hyperparameter has the greatest effect on a

model's accuracy

to determine which hyperparameter has the least effect on a model's accuracy

✓ to apply different combinations of hyperparameter settings to see if test set performance can be improved

✗ to apply different combinations of hyperparameter settings to see if the training performance can be improved

✓ Question 12 of 36

The function that takes a sentence and splits it in a string of words is _____ .

tkinter

textcat

✓ tokenize

pos_tag

✓ Question 13 of 36

What is true about TF-IDF?

✓ all of these answers

The cells represent a weighting,

There is still one row per text message.

The columns still represent single unique terms.

✓ Question 14 of 36

NLP is a field concerned with the ability of a computer to _____ human language.

manipulate

understand

analyze

✓ all of these answers

✓ Question 15 of 36

What does an n-gram represent?

all the words of length `n` in a string

all the sentences of length `n` inside a data file

all the words that appear `n` times inside a data file

✓ all combinations of `n` adjacent words

✓ Question 16 of 36

In bigram, how many tokens will be generated for the statement "NLP is an interesting topic"?

5

✓ 4

2

3

✓ Question 17 of 36

Why is lemmatizing more accurate in finding word variations that have the same meaning?

✓ It has a database of nouns, verbs, adjectives, and adverbs that are grouped together as sets of synonyms.

It uses context clues from other words in the sentence.

It has a centralized database of nouns, verbs, adjectives, and adverbs that are constantly updated by its userbase.

It specializes in finding synonyms for common slang words found in SMS messages.

✓ Question 18 of 36

Lemmatizing 'meanness' and 'meaning' results in _____ .

✓ meanness
meaning

mean
mean

meaning
meanness

meanness
meanness

✓ Question 19 of 36

Random forest is an example of what type of method?

✓ ensemble

vectorization

boosting

matrix

Question 20 of 36

Why do you need a holdout test set?

to ensure that the remaining data has an equal sample size

to reduce the amount of data subsets we need to evaluate

✓ to evaluate a model's ability to generalize to unseen data

Question 21 of 36

Which of these is NOT an example of supervised learning?

determining what height percentile a child is based on his or her current age

✓ grouping together similar emails into distinct folders based on the content

determining which emails are spam based on known information

about the sender

Question 22 of 36

How does the Box-Cox power transformation work?

It provides the equation of a line that has the best correlation with the actual data points.

It applies a range of logarithms to your data points to determine if the result fits closely to a normal distribution.

✓ It applies a range of exponents to your data points to determine if the result fits closely to a normal distribution.

It provides a mapping of the actual data points to data points that would be found in a normal distribution.

Question 23 of 36

If X is -2, which one of the following is a correct transformation?

✓ 1/y^2

x^-2

x^2

y^2

Question 24 of 36

Which statement best describes a bimodal distribution graph?

a curve with one spike and two long tails on each side of the spike

a curve with two large dips in different locations

✓ a curve with two large spikes in different locations

a curve with a long tail to the right

Question 25 of 36

Lemmatizing _____ .

is faster than stemming

chops the end of the word using heuristics

✓ is slower than stemming

does not understand the context in which the word is used

Question 26 of 36

Why is it necessary to analyze the output of a stemming process?

- ✓ **Stemming algorithms are not perfect, as they can stumble on slang words and certain root words.**
- **Stemming algorithms can inadvertently remove certain tokens.**
- **Stemming algorithms can crash on long words and possibly return incomplete output.**

Question 27 of 36

What is the result of running the stemmer against 'run', 'running', and 'runner'?

- ✓ **run**
  run
  runner
- **runner**
  run
  runner
- **running**
  running
  runner

run
run
run

Question 28 of 36

Why is stemming important?

✗ It captures variations of the same root word to help an NLP algorithm learn more words.

It captures variations of the same root word to produce a larger number of tokens.

It indexes all variations of a word to populate a NLP word database.

✓ It reduces variations of the same root word to produce a smaller number of tokens.

Question 29 of 36

Stemming Meanness/meaning will result in _____ .

Meanness

✓ Mean

Mea

meaning

Question 30 of 36

Using the NLTK, which package can be called to get a pre-defined list of stop words?

nltk.stopwords.words

✓ nltk.corpus.stopwords.words

nltk.corpus.words.stopwords

nltk.stopwords

Question 31 of 36

What will be displayed on the screen when the following code runs?

```
tokenized = ['test','in','the','rest','of','for','new','last']
result = [word for word in tokenized if word not in
['in','on','the','of','for']]
print(result)
```

['test', 'rest', 'new', 'for']

✓ ['test', 'rest', 'new', 'last']

```
['test', 'for', 'the', 'rest', 'new', 'last']
```

```
['test', 'new', 'last']
```

Question 32 of 36

When writing your own tokenization function, which essential step must your function be able to do?

Check that your function doesn't read past the end of the data file.

Check that the correct punctuation is used for any sentences found in the data file.

data file.

Return the number of tokens found in the data file.

Question 33 of 36

What will be printed with the following statement:

```
print(re.split('\W+',"some of the-words are+combined"))
```

✓ `['some', 'of', 'the', 'words', 'are', 'combined']`

```
['some', 'of', 'the-words', 'are+combined']
```

```
['', ' ', ' ', '-', ' ', '+', '']
```

```
['some of the', 'words are', 'combined']
```

Question 34 of 36

What stop word that will be removed from this sentence: "This is a test for the man to be successful in their lives"?

- the
- to
- for
- ✓ all of these answers

Question 35 of 36

In the spam-ham example, what does a recall rate of 55.2% mean?

- ✓ 55.2% of spam properly went to the spam folder while the rest went to the inbox.
- 55.2% of all emails were identified as spam.

Any email, whether it was spam or not, was correctly identified 55.2% of the time.

Question 36 of 36

What does it mean when a random forest has `max_depth=none` and `n_estimators=10`?

the random forest will have 10 decision trees with a minimum depth of zero

the random forest will have unlimited trees with a depth of 10

✓ the random forest will have 10 decision trees of unlimited depth

Return to course